

Project: From Model to Production: Scalable Batch Processing for Sustainable Retail

DLBDSMTP01 Task 2: Batch Image Classification for Refund Department

Alejandro Moral Aranda

IU University

August 24, 2025

Try it online now: <https://api-fenlei-production.up.railway.app/>
Hosted on Railway (free plan). First request may be slow after inactivity.
Also runnable locally.
Source Code: github.com/alemxral/api-fenlei

Web Interface (Screenshot)

♥ API Status

🟢 API is healthy - MobileNetV2 model loaded



Image Classification



Prediction History



Single Image Classification

Select Image

Choose File

plane.jpeg

Supported formats: JPG, PNG, GIF, BMP, WEBP (max 16MB)

Number of Predictions

Top 5



Classify Image

Preview:



Batch Image Classification

Select Images

Choose Files

No file chosen

Select up to 10 images (max 16MB each)

Number of Predictions per Image

Top 5



Classify Images

Outline

- 1 Project Overview
- 2 Preparations & Planning
- 3 Architecture
- 4 Challenges & Requirements
- 5 Risks, Timeline & Resources
- 6 Results & Reflection
- 7 Summary & Conclusion

Goal: Automate refund item categorization from images to reduce manual effort and costs.

Solution: REST API for batch image classification using MobileNetV2.

Deployment: Hosted on Railway (public demo), also supports local runs.

Key Features:

- Batch and single image processing.

- Web interface for interactive/manual testing.

- Structured API responses and validation.

Objectives

Reduce manual sorting workload.

Enable nightly batch classification jobs.

Ensure robust data processing and validation.

Provide easy-to-use manual testing UI.

Preparations

Model: Pre-trained MobileNetV2 (ImageNet).

Environment: Python 3.11, TensorFlow/Keras, Flask.

Validation: File type, size, dimensions, integrity checks.

Prototyping: Verified local inference and batch processing.

Architecture: Defined ingestion, storage, processing, and API flows.

Code Structure: Modular organization (API, ML logic, utilities, UI).

Deployment: Public demo on Railway, portable for local use.

Why This Approach?

Pre-trained Model: Fast, high-quality baseline.

MobileNetV2: Efficient CPU inference.

Batch Processing: Designed for nightly automation.

Flask: Simple, reliable RESTful backend.

Railway: Easy/cheap cloud demo deployment.

System Architecture

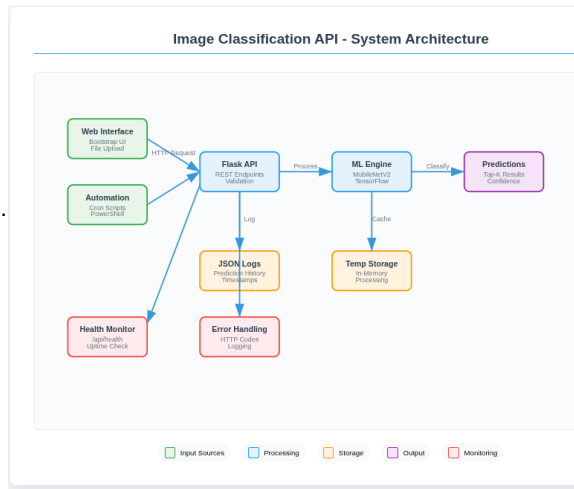
Frontend: HTML/CSS, Vanilla JS for UI and API interaction.

Backend: Flask REST API.

ML Model: MobileNetV2.

Batch: Scheduled jobs (cron, GitHub Actions).

Storage: JSON logs, optional DB backend.



System architecture diagram

API Endpoints

POST /api/classify Classify a single image.

POST /api/classify/batch Batch classification (up to 10 images).

GET /api/health API health check.

GET /api/info Build/config info.

Clear JSON errors for invalid format, size, or corrupted images.

Checks for image dimensions (324096px) and MIME/type.

Consistent and meaningful HTTP status codes.

Batch Processing

Nightly jobs for batch image classification.

Triggered by cron or GitHub Actions.

Logs written to temporary storage or database.

Integration Challenges

Speed Needs:

Batch: high throughput.

Interactive: low latency.

Scalability:

Handle more users/data via horizontal scaling.

Consistent Preprocessing:

Same data pipeline for training and serving.

Security & Privacy:

Encryption, access control.

Resource Limits:

Stay within budget/hardware constraints.

Model Versioning:

Track and roll back deployed models.

Data Requirements

Getting Images: Collect photos easily from warehouse devices.

Where to Store:

- Save images in object storage (like S3 or minIO).

- Keep results in JSON files or a database.

- Save extra info (metadata) in a database (like PostgreSQL).

Processing Steps:

- Check and label images for quality.

- Always use the same cleaning and prep steps.

How It Works: Special programs check, organize, and save every image and its info.

Data Quality: Resolution, corruption, metadata checks.

Prediction Drift: Monitor class distribution vs. baseline.

Performance: Latency, throughput, error rates.

Model Quality: Periodic accuracy/F1 checks.

Alerting: Threshold-based notifications.

Audit Logs: Track prediction events and model versions.



Image Classification API - System Workflow



Overnight Automation Workflow

Linux/macOS

- image_classifier_automation.sh
- Cron: 0 2 * * *
- Bash scripting
- Comprehensive logging

Windows

- image_classifier_automation.ps1
- Task Scheduler
- PowerShell scripting
- Progress tracking

Health Check

Validate Images

Batch Process

Log Results

Generate Reports



Input Sources



Processing



Storage



Output



Monitoring

Infra: Cold starts, resource limits on free hosting.

Compatibility: TensorFlow/Pillow binary dependencies.

Security: File upload validation, size caps.

Performance: CPU-only inference latency.

Mitigations: Health checks, format validation, worker tuning.

Compute: Railway free tier, local CPU for development/testing.

Software: Open-source stack (Flask, TensorFlow/Keras, Pillow).

Delivered REST API and batch processing for refund categorization.

Public demo: <https://api-fenlei-production.up.railway.app/>.

Met objectives for functionality and usability.

Process: Iterative development for fast feedback.

Resources: Free-tier hosting is viable for demos; scalability remains to be tested.

Lessons: Early validation and health endpoints reduce downtime.

Summary & Conclusion

Developed a scalable API to automatically classify refund item images, reducing manual effort and errors.

Supports both real-time and batch processing, making it flexible for different needs.

Modular design allows easy updates, integration, and future improvements.

The project demonstrates how AI can streamline retail workflows and can be extended to new use cases.

API Usage: curl Examples

Health check

```
curl https://api-fenlei-production.up.railway.app/api/health
```

Single image classification (local)

```
curl -X POST -F "image=@path/to/image.jpg" -F "top_k=5" \  
http://localhost:5000/api/classify
```

API information

```
curl http://localhost:5000/api/info
```

Bibliography



Mark Sandler, Andrew Howard, Menglong Zhu, Andrey Zhmoginov, Liang-Chieh Chen.
MobileNetV2: Inverted Residuals and Linear Bottlenecks.
[In CVPR, 2018.](#)



Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, Li Fei-Fei.
ImageNet: A Large-Scale Hierarchical Image Database.
[In CVPR, 2009.](#)



Martín Abadi et al.
TensorFlow: Large-Scale Machine Learning on Heterogeneous Systems, 2015.
<https://www.tensorflow.org/>



François Chollet et al.
Keras, 2015.
<https://keras.io/>



Armin Ronacher.a
Flask (Python Web Framework).
<https://flask.palletsprojects.com/>



Questions?

Live Demo: <https://api-fenlei-production.up.railway.app/>

Source Code: github.com/alemxral/api-fenlei